

DOCKER + KUBERNETES

Teaching Faculty: Umur INAN

Prepared by Umur INAN

DOCKER

- Docker is an open platform for developing, shipping, and running applications.
- Docker enables you to separate your applications from your infrastructure so you can deliver software quickly.
- Docker provides the ability to package and run an application in a loosely isolated environment called a container.
- Containers are lightweight and contain everything needed to run the application, so you do not need to rely on what is currently installed on the host.

CONTAINER

- Container is a software package that consists of all the dependencies required to run an application.
- A Docker container image is a lightweight, standalone, executable package of software that includes everything needed to run an application: code, runtime, system tools, system libraries and settings.
- Containerization is simply a way of packaging software applications into containers.

VIRTUALIZATION

- Virtualization is the process of creating a simulated computing environment that's abstracted from the physical computing hardware.
- Virtualization allows you to create multiple, virtual computing instances from the hardware and software components of a single machine.

HYPERVISOR


- The software that enables virtualization is called a hypervisor.
- It's a lightweight software layer that sits between the physical hardware and the virtualized environments and allows multiple operating systems (OS) to run in tandem on the same hardware.
- The hypervisor is the middleman that pulls resources from the raw materials of your infrastructure and directs them to the various computing instances.

VIRTUAL MACHINES


- Separate computers running on hardware that is contained in one physical computer.
- Each VM requires its own OS.
- The OS and any applications running on an individual VM share hardware resources from a single host server, or from a pool of host servers.

DOCKER VS VIRTUALIZATION

Docker Containers



Virtual Machines



BUILDING DOCKER IMAGES

- Docker image is built using a Docker file.
- A docker image contains all the project's code, whereas a Dockerfile is a text file which contains commands for building a Docker image.
- With docker images containers can be built.
- From Docker Hub users can pull any Docker Image and build new containers.

DOCKER HUB

- Docker Hub is a hosted repository service provided by Docker for finding and sharing container images.
- Users get access to free public repositories for storing and sharing images or can choose subscription plan for private repos.

DOCKER COMPOSE

- Docker Compose is used for running multiple containers as a single service.
- Compose is a tool for defining and running multi-container Docker applications.
- With Compose, you use a YAML file to configure your application's services.

DOCKER COMPOSE

- Then, with a single command, you create and start all the services from your configuration.
- Docker Compose works by applying many rules declared within a single `docker-compose.yml` configuration file.

DOCKER-COMPOSE.YML

- Contains
 - services
 - volumes (optional)
 - networks



```
docker-compose up
```

VOLUMES + NETWORKS

- Volume
 - is a shared directory in the host, visible from a container.
- Network
 - define the communication rules between containers, and between a container and the host.
 - Common network zones will make containers' services discoverable by each other, while private zones will segregate them in virtual sandboxes.

SERVICES

- Refers to containers' configuration.
- Let's dockerize an app consisting of frontend-app, backend-app, database and app-umur.

services:

frontend-app:

image: my-frontend-app

ports:

- "8080:2000"

backend-app:

image: my-backend-app

db:

image: postgres

umur-app:

image: awesome-app

Port 2000 will be available in 8080 on the host.

SERVICES

```
services:  
  frontend-app:  
    image: my-frontend-app  
    ports:  
      - "8080:2000"  
    networks:  
      - network1  
  backend-app:  
    image: my-backend-app  
    networks:  
      - network1  
  db:  
    image: postgres  
  umur-app:  
    image: awesome-app  
    networks:  
      - network2
```

```
networks:  
  network1: {}  
  network2: {}
```

my-backend-app can reach frontend-app
whereas it cannot reach umur-app.

VOLUMES

```
services:
  frontend-app:
    image: my-frontend-app
    ports:
      - "8080:2000"
    networks:
      - network1
    volumes:
      - /tmp:/volume1/from-host
  backend-app:
    image: my-backend-app
    networks:
      - network1
```

```
  db:
    image: postgres
  umur-app:
    image: awesome-app
    networks:
      - network2
```

```
networks:
  network1: {}
  network2: {}
```

The /tmp folder of the host is mapped to /volume1/from-host folder of the container.

DOCKER COMPOSE COMMANDS

- `docker compose up`
 - Creates and start containers.
- `docker compose down`
 - Stops and removes containers.
- `docker compose create`
 - Creates containers for a service.
- `docker compose ls`
 - Lists running compose projects.

DOCKER COMMANDS

- `docker ps`
 - Lists the running containers.
- `docker ps -a`
 - Lists all the running and exited containers.
- `docker stop <container id>`
 - Stops a running container
- `docker images`
 - Lists all the locally stored docker images.

DOCKER COMMANDS

- `docker rm <container id>`
 - Deletes a stopped container.
- `docker build <path to docker file>`
 - Builds an image from a specified docker file.
- `docker run -pHOST:CONTAINER -d --name custom-name -net network-name name-of-image`

DOCKER COMMANDS

- `docker logs container-id`
- `docker pull name-of-image`
- `docker exec -it container-id /bin/bash`
- `docker start name-of-image`

DOCKER COMMANDS

- `docker network ls`
- `docker network create name-of-network`
- `docker exec -it container-id /bin/bash`
- `docker start name-of-image`

DOCKERFILE

- Each Dockerfile is a script, composed of various commands (instructions) and arguments listed successively to automatically perform actions on a base image in order to create (or form) a new one.
- Dockerfiles begin with defining an image FROM which the build process starts. Followed by various other methods, commands, and arguments (or conditions), in return, provide a new image which is to be used for creating docker containers.



```
docker build
```

DOCKERFILE COMMANDS

- MAINTAINER
 - Sets the author.
 - MAINTAINER umur_inan
- ADD
 - Takes 2 arguments, a source and a destination.
 - It copies files from source on the host machine into the container.
 - The source can also be a URL.
 - ADD /from/host /to/container

DOCKERFILE COMMANDS

- CMD
 - Executes a specific command.
 - It is not executed during build but when the container is created/started.
 - CMD "echo" "Hello World"
- RUN
 - It is executed during build unlike CMD command.

DOCKERFILE COMMANDS

- FROM
 - Specifies the base image to use to create the container.
 - It can be default images from docker hub or your own custom image.
 - FROM <image>
 - FROM <image>:<tag>
 - Tag is latest by default.

DOCKERFILE COMMANDS

- WORKDIR
 - Sets the current path or the path where CMD, RUN, COPY, ADD commands going to be executed.
 - WORKDIR /source
- VOLUME
 - Creates a mount point with the specified name and marks it as holding externally mounted volumes from native host or other containers.

DOCKERFILE COMMANDS

- EXPOSE
 - EXPOSE is used as documentation for the port. This is just a communication between the person who builds the image and the person who runs the container.
 - EXPOSE 8080

DOCKERFILE COMMANDS

- `docker images`
 - List all images.
- `docker run -it -d image-name`
 - Run the image.
- `docker build -t image-name:tag-name .`
 - Builds a docker image.

DOCKERFILE

```
# Comment Line
FROM ubuntu
MAINTAINER uinan@miu.edu

WORKDIR /home/umurinan

RUN apt-get update
RUN apt-get install -y nginx

CMD ["echo", "Hello World"]

COPY from-my-computer to-container
```

SPRING BOOT -- DOCKER

```
FROM openjdk
```

```
WORKDIR /path/to/jar/file
```

```
COPY app.jar app.jar
```

```
EXPOSE 8080
```

```
CMD ["java", "-jar", "app.jar"]
```

KUBERNETES

- Kubernetes, also known as K8s, is an open-source system for
 - automating deployment,
 - scaling,
 - management of
- containerized applications.
- It groups containers that make up an application into logical units for easy management and discovery.

CONTAINER ORCHESTRATOR

- Container orchestrators are designed to run complex applications with large numbers of scalable components.
- They work by inspecting the underlying infrastructure and determining the best server to run each container.

KUBERNETES

- Open-source
 - You can download and use it without paying any fee.
- Battle-tested
 - There're plenty of examples of companies running it in production.
- Well-looked-after
 - Redhat, Google, Microsoft, IBM, Cisco are only a few of the companies that have heavily invested in the future of Kubernetes by creating managed services, contributing to upstream development and offering training and consulting.

KUBERNETES FEATURES

- Service Discovery and Load Balancing
- Storage Orchestration
- Automated Rollouts and Rollbacks
- Automatic Bin Packing
- Self-Healing
- Secret and Configuration Management

SERVICE DISCOVERY AND LOAD BALANCING

- Kubernetes can expose a container using the DNS name or using their own IP address.
- If traffic to a container is high, Kubernetes is able to load balance and distribute the network traffic so that the deployment is stable.

STORAGE ORCHESTRATION

- Kubernetes allows you to automatically mount a storage system of your choice, such as local storages, public cloud providers, and more.

AUTOMATED ROLLOUTS AND ROLLBACKS

- Kubernetes allows you to automatically mount a storage system of your choice, such as local storages, public cloud providers, and more.

AUTOMATIC BIN PACKING

- You provide Kubernetes with a cluster of nodes that it can use to run containerized tasks.
- You tell Kubernetes how much CPU and memory (RAM) each container needs.
- Kubernetes can fit containers onto your nodes to make the best use of your resources.

SELF-HEALING

- Kubernetes restarts containers that fail, replaces containers, kills containers that don't respond to your user-defined health check, and doesn't advertise them to clients until they are ready to serve.

SECRET AND CONFIGURATION MANAGEMENT

- Kubernetes lets you store and manage sensitive information, such as passwords, OAuth tokens, and SSH keys.
- You can deploy and update secrets and application configuration without rebuilding your container images, and without exposing secrets in your stack configuration.

TERMINOLOGY

- Namespaces
 - In Kubernetes, the namespaces is effectively your working area. It's like a project in GCP or a similar thing in AWS.

TERMINOLOGY

- Pods
 - Abstraction for Containers.
 - Pod is a group of containers and is the smallest unit that Kubernetes administers.
 - Pods have a single IP address that is applied to every container within the pod.
 - Containers in a pod share the same resources such as memory and storage.
 - It's quite common to have a pod with only a single container, when the application or service is a single process that needs to run.

TERMINOLOGY

- Service
 - A service is an abstraction over the pods, and essentially, the only interface the various application consumers interact with.
 - As pods are replaced, their internal names and IPs might change.
 - A service exposes a single machine name or IP address mapped to pods whose underlying names and numbers are unreliable.
 - A service ensures that, to the outside network, everything appears to be unchanged.

TERMINOLOGY

- Service
 - Has permanent IP address.
 - Is also a load balancer.

TERMINOLOGY

- Deployment
 - Blueprint for pods.
 - Specify how many replicas you want to have
 - Abstraction of Pods.
 - For stateLESS applications.
 - Use statefulset for stateful applications or databases.

TERMINOLOGY

- Ingress
 - This works with the service to make sure everything ends up in the right place. Ingress can also provide load balancing.
- Node
 - A Kubernetes node manages and runs pods; it's the machine (whether virtualized or physical) that performs the given work.
 - Just as pods collect individual containers that operate together, a node collects entire pods that function together.

TERMINOLOGY

- API Server
 - The API server exposes a REST interface to the Kubernetes cluster.
 - All operations against pods, services, and so forth, are executed programmatically by communicating with the endpoints provided by it.

TERMINOLOGY

- Scheduler
 - The scheduler is responsible for assigning work to the various nodes.
 - It keeps watch over the resource capacity and ensures that a worker node's performance is within an appropriate threshold.
- Controller Manager
 - The controller-manager is responsible for making sure that the shared state of the cluster is operating as expected.
 - More accurately, the controller manager oversees various controllers which respond to events

TERMINOLOGY

- Kubelet
 - A Kubelet tracks the state of a pod to ensure that all the containers are running.
 - It provides a heartbeat message every few seconds to the control plane.
 - If a replication controller does not receive that message, the node is marked as unhealthy.

TERMINOLOGY

- Kube Proxy
 - The Kube proxy routes traffic coming into a node from the service.
 - It forwards requests for work to the correct containers.
- Etcd
 - Holds the current status of any K8s component.
 - is a distributed key-value store that Kubernetes uses to share information about the overall state of a cluster.

CONFIGURATION FILES

- Each config file has 3 parts.
 - Metadata
 - Specification
 - Attributes of spec are specific to the kind.
 - Kind: Deployment, Service
 - Status
 - Auto-generated by Kubernetes.
 - Desired state != Actual State
 - It comes from etcd.

COMMANDS

- `kubectl get node`
- `kubectl apply -f config_file_path`
- `kubectl get all`
- `kubectl describe service name-of-the-service`
- `kubectl describe pod name-of-the-pod`
- `minikube ip`
- `kubectl port-forward svc/webapp-service 8080:8080`

KUBERNETES ARCHITECTURE

