



Software Architecture

Service Discovery



What do you know about
Service Discovery?



Service discovery

Back Ends

- Stand alone registry
 - Spring supported :
 - Consul (+ registrator for containers)
 - Eureka
 - ZooKeeper, Etcd, ...
- Platform
 - Kubernetes service catalog
 - DNS – service LB based
 - Istio “Service mesh”
 - Control Plane

Feature Overlap!

Capabilities				Spring Cloud with Kubernetes & Istio on IaaS+
DevOps Experience				Self service, multi-environment capabilities
Auto Scaling & Self Healing				Pod/Cluster Autoscaler, HealthIndicator, Scheduler, Pool Ejection
Deployment & Scheduling				Deployment strategy, DarkLaunch, A/B, Canary, Scheduler Strategy
Resilience & Fault Tolerance				HealthIndicator, Hystrix, HealthCheck, Process Check, Circuit Breaker/Timeout/Retry
Distributed Tracing				Zipkin, Jaeger
Centralized Metrics				Heapster, Prometheus, Grafana
Centralized Logging				EFK
API Gateway				Zuul, Traffic Control, Egress
Load Balancing				Ribbon, Service, Envoy
Chaos Engineering				Chaos Monkey for Spring Boot, Chaos Toolkit Kubernetes, Envoy
Service Discovery				Service
Configuration Management				Externalized Configuration, ConfigMaps, Secrets
Application Packaging				Spring Boot Maven/Gradle plugin
Job Management				Spring Batch, Scheduled Jobs
Process Isolation				Docker, Pods, Envoy
Environment Management				Namespaces, Authorization
Resource Management				CPU and Memory Limits, Namespace resource quotas
Operating System				Ubuntu, Atomic, ...
Virtualization				VMWare, Openstack, ...
Hardware, Storage, Networking				AWS, GCP, Azure, ...

<https://mrumpf.github.io/spring-cloud-on-k8s/>



OK, back to the basics

- What's container orchestration
 - What's Kubernetes
 - <https://www.youtube.com/watch?v=R-3dfURb2hA>



What's a service mesh

- Control plane , add proxy/side car to a service
 - Communication now goes through the proxy
 - Define policies
 - Routing, security/ACL , traffic management, telemetry , even service discovery,...



Show me how?

- Building cloud-native applications with Kubernetes and Istio, by Kelsey Hightower
 - <https://www.youtube.com/watch?v=6BYq6hNhcel>

Architecture implications

- Moving cross-cutting logic from services → Implement at platform level (k8s+istio)
 - No need to integrate with a stand alone discovery backend (consul,eureka ,..)
 - Remove circuit breaking,retries, client LB logic from services
 - Telemetry, ACL ,.. : dynamic and no need to change services “almost free”
 - Flexibility: not just java services or framework with integration with certain back ends



Architecture implications

- Cons:
 - Complexity , learning curve
 - More dependence on the deployment environment → less flexibility to change later



Still, we need DIY

- e.g. Deploy a dedicated Registry (Consul) cluster on top of platform (e.g. k8s)
 - Works best for legacy/Java Microservices
 - Offers more flexibility → no lock in platform
 - Cons:
 - Apps need to be configured
- Check this out:
 - <https://luizkowalski.net/deploying-spring-cloud-netflix-apps-on-kubernetes/>



OK, Let's see a sample service registry
“Consul”



Confused?

- Check Consul vs other:
 - <https://www.consul.io/intro/vs/index.html>
- Checkout:
 - <https://attx-project.github.io/Service-Discovery-Solutions.html>



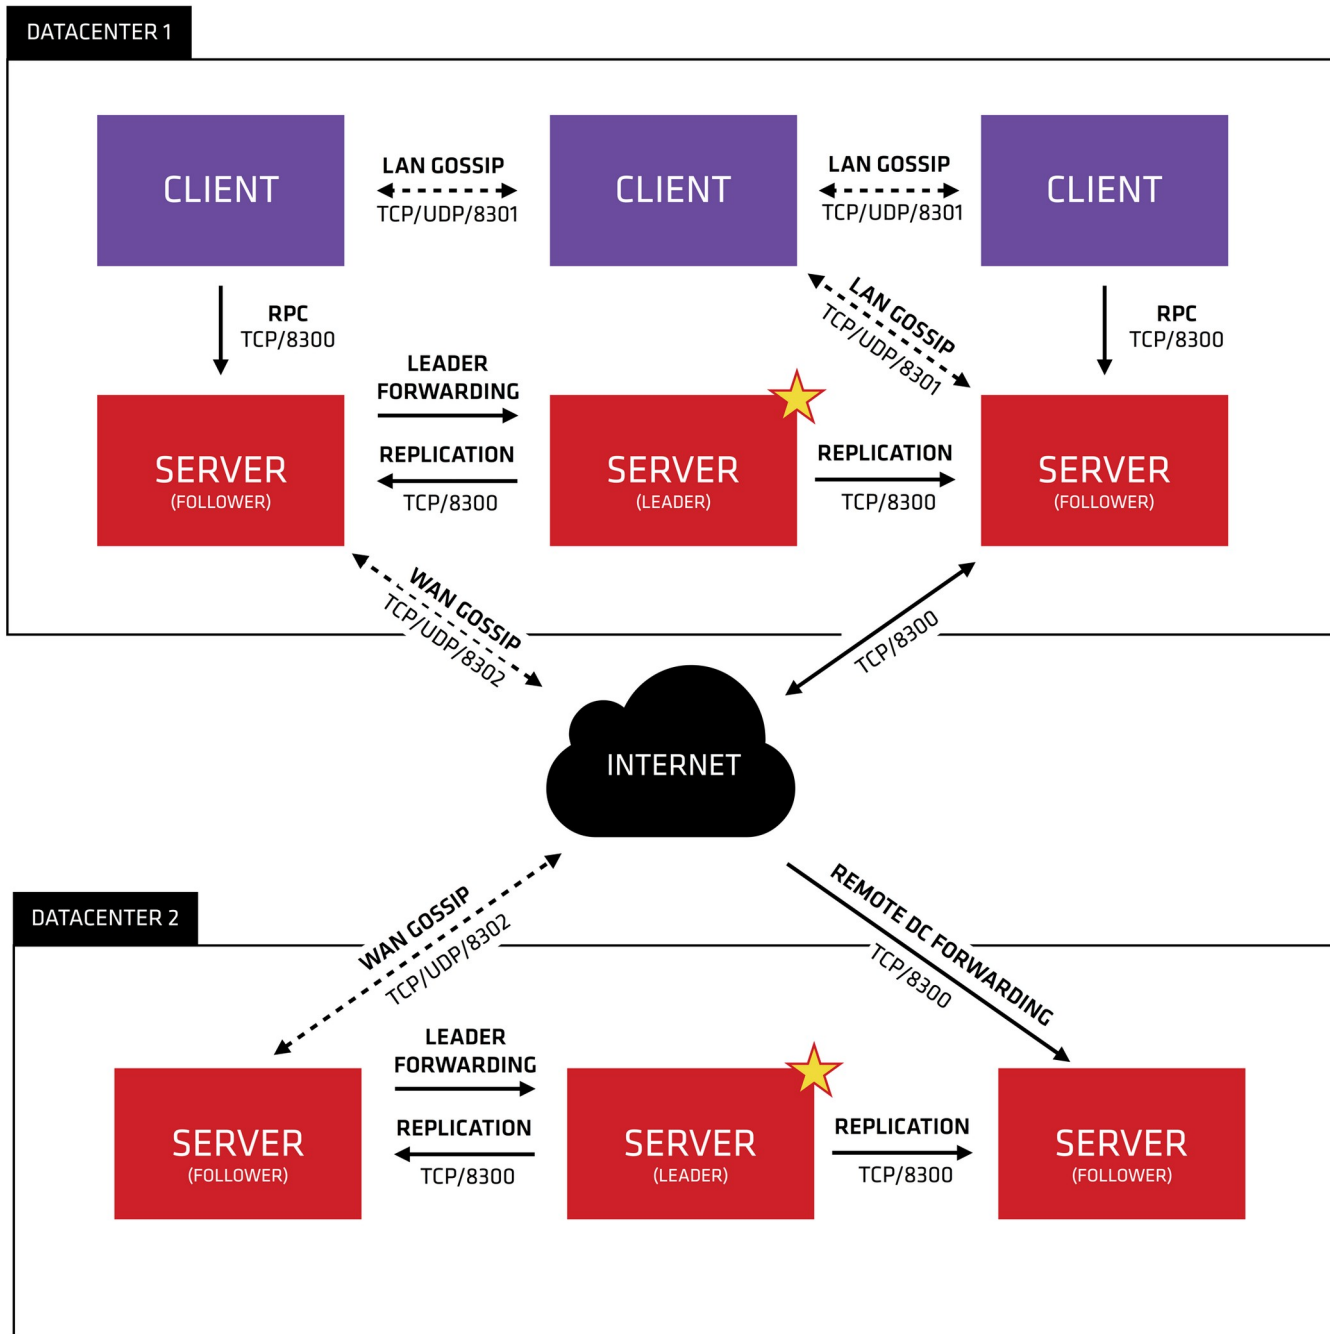
Service Discovery

Use case : Consul



Why we need it

- Let's hear the creator
 - Introduction to HashiCorp Consul | Armon Dadgar
 - <https://www.youtube.com/watch?v=mxzMdl0KvBI>





What's going on?

- How nodes inside DC find each other?
- How leader election works?
- How followers consent on the values?
- How nodes between 2 Data centers connect?
- What's stored in each DC?
 - Hint: Data is different, but request can be forwarded



how a node knows about other nodes ?



But how a node knows about other nodes ?

Gossiping!



Gossip protocol

- Also known as epidemic protocol
- A family of protocol styles and implementations
- Consul is based on serf library which implements **SWIM** protocol
 - **S**calable **W**eakly-consistent **I**nfection-style **P**rocess Group **M**embership Protocol



SWIM Gossiping

- Got plenty of time, read the paper:
 - http://www.cs.cornell.edu/info/projects/spinglass/public_pdfs/swim.pdf



Consul Gossiping

- LAN:
 - All Nodes participate (server and clients)
- WAN:
 - Across Internet between data centers
 - Only servers participate



How they elect a master and consent on values?



How they elect a master and consent on values?

Consensus Protocol



Raft

- In CAP terms : Strong consistency (CP)
- Solves distributed consensus problem
- Don't expect high throughput
 - Implemented in systems like
 - ZooKeeper , Consul , chubby ...



Raft

- Let's see how it works
 - <http://thesecretlivesofdata.com/>



Raft In consul

- Only server nodes keeps state (not clients)
- Limit # servers
 - Raft communication is noisy
 - Typically odd # server , e.g. 3,5 to maintain majority



Distributed configuration

- Consul
- Spring cloud config
 - Default storage is git, among many other
- Etcd
- ...



Final Thoughts

- Now we can register services and discover them
 - By HTTP or DNS (SRV records)
- Monitor health checks!
- Also store K/V for distributed configuration!
 - Treat profiles as hierarchy